



SEO STRATEGY DOC

Search + AI Discovery Operating System

An evidence-based strategy for approximately 15 service, SaaS, local, ecommerce, and content-led websites.

Version 2.0 | Evidence review 2026-07-17

Executive Strategy

The objective is not to chase a universal checklist. It is to create the strongest defensible system for being discovered, selected, trusted, cited, and converted across search and AI-assisted experiences.

Strategic position

Search and AI visibility share the same foundation: accessible and indexable pages, unambiguous entities, original and supportable information, strong task completion, and reliable technical delivery. No legitimate system can guarantee a #1 ranking or citation; this strategy maximizes controllable quality and builds measurement into every decision.

The six operating principles

1. Intent before format	Page depth and structure follow the user's task. No universal word, keyword, image, heading, FAQ, or link quota.
2. Truth before scale	No invented authors, credentials, dates, reviews, locations, case studies, metrics, or schema properties.
3. Independent domain value	Every portfolio site needs a distinct entity, audience, purpose, and original value. Token-swapped pages and doorway networks are prohibited.
4. Page-specific controls	Indexation, content, schema, links, freshness, and conversion rules differ by page type and business model.
5. Evidence hierarchy	Engine requirements and web standards outrank vendor claims. Recommendations, experiments, and internal conventions are labeled.
6. Outcomes with guardrails	Measure discovery, visibility, qualified engagement, revenue, and risk. Experiments require hypotheses and rollback conditions.

P0 corrections made

- Removed deceptive publication backdating and false freshness.
- Prohibited fabricated author identities, credentials, customer evidence, and outcomes.
- Replaced universal schema with truthful page-entity mapping.
- Removed forced .gov/.edu links, keyword density, and arbitrary content quotas.
- Reframed llms.txt as optional experimentation, not AI citation control.
- Replaced ranking guarantees with measurable, evidence-based controls.

Operating Model

1. Site charter

Define business model, audience, markets, entity, products/services, conversion, eligibility, YMYL exposure, evidence assets, and unique value before producing pages.

2. Query-to-page map

Assign each material intent to one preferred URL. Split only for a real difference in audience, offer, location, language, or desired result. Audit cannibalization before expansion.

3. Page-type contract

Every URL gets an explicit indexation state, purpose, main entity, unique value, metadata, internal-link role, schema decision, freshness trigger, owner, and acceptance test.

4. Evidence-led production

Brief, draft, evidence review, experience review, SEO/accessibility review, approval, publish validation, and monitoring form one accountable workflow.

5. Technical reliability

Align status codes, canonicals, robots, sitemaps, links, rendered HTML, hreflang, facets, pagination, structured data, performance, security, and accessibility.

6. Portfolio governance

Maintain a registry across all domains. Reuse code and controls; do not clone value propositions or create a manipulative cross-link network.

Rule classification

REQUIRED	Technical validity, accessibility, compliance, truth, or intended indexation
RECOMMENDED	Evidence-based default; document deviations
CONDITIONAL	Applies to a defined page, market, feature, or risk
OPTIONAL	Useful in context; not an SEO requirement
EXPERIMENTAL	Measured hypothesis with a stop condition
PROHIBITED	Deceptive, spam-oriented, fabricated, inaccessible, or materially risky

Reusable Prompt - Part A

Character count: 2,010 (limit: 4,000)

PART A - RESEARCH AND AUDIT

Act as a senior technical SEO, content-strategy, local/ecommerce SEO, accessibility, and AI-discovery auditor. Audit the current playbook at <https://github.com/fabiodemelo/webdev/blob/main/playbooks/content-seo/SPEC.md> for use across approximately 15 independently valuable service, SaaS, local, ecommerce, and publishing sites.

Use current primary sources first: Google Search Central/Search Essentials/spam policies, Bing Webmaster guidance, Schema.org plus engine-specific rich-result rules, Merchant Center, Google Business Profile, Core Web Vitals, and WCAG 2.2. Cite every material recommendation. Label its evidence as ENGINE REQUIREMENT, ENGINE GUIDANCE, WEB STANDARD, INDUSTRY PRACTICE, INTERNAL CONTROL, or EXPERIMENT. Never treat correlation as a confirmed ranking factor; flag uncertainty and date-sensitive guidance.

Audit every existing rule. Challenge keyword counts/density, fixed word or heading counts, link/image/FAQ/CTA quotas, title/meta ranges, authorship and E-E-A-T claims, forced .gov/.edu links, FAQ schema, global schema, llms.txt, crawler rules, backdating, bulk launch content, AI-generated material, competitor claims, Core Web Vitals, canonical/hreflang/pagination/facets, and ranking/conversion promises. Identify spam, scaled-content, doorway, duplicate, trust, legal, accessibility, and reputation risks.

Classify every rule: REQUIRED, RECOMMENDED, CONDITIONAL, OPTIONAL, EXPERIMENTAL, or PROHIBITED. Produce: (1) executive assessment; (2) findings matrix with current rule, evidence, risk, Keep/Revise/Remove/Replace/Add, and P0-P3 priority; (3) critical risks; (4) missing systems; and (5) recommended document architecture.

Do not promise rankings. Optimize for the strongest defensible foundation: crawlability, indexability, intent satisfaction, original value, truthful entities, accessibility, performance, trust, qualified conversion, and measurable improvement. Stop after the audit and architecture. Preserve the findings for Part B.

Reusable Prompt - Part B

Character count: 2,254 (limit: 4,000)

PART B - REBUILD, VALIDATE, AND UPDATE

Continue from Part A. Rewrite the playbook as a copy-ready operating specification and update SPEC.md. Preserve good rules, replace unsupported formulas, remove risky tactics, and add missing systems. Every normative rule must use REQUIRED, RECOMMENDED, CONDITIONAL, OPTIONAL, EXPERIMENTAL, or PROHIBITED and, where material, its evidence type.

Create separate standards for: homepage; service; SaaS feature/product landing; ecommerce product; category/collection; pricing; industry/audience; location/local service; article; pillar; tutorial; research/data; case study; comparison/alternative; about/contact; author/expert; help/FAQ; legal/security; search/filter/tag/archive/pagination; campaign landing; and login/account/checkout/admin. For each define intent, indexation, unique value, content, metadata/H1/URL/canonical/robots/hreflang, links/citations, media, conversion, trust, schema, accessibility, freshness, acceptance checks, and anti-patterns. Use ranges only as non-ranking starting points when supported.

Add systems for query-to-page mapping, architecture, topic hubs, cannibalization, entity consistency, briefs, inventories/gaps, consolidation/redirects/migrations, crawl/index/render governance, JavaScript, duplicates, facets/pagination, international/local/ecommerce/image/video SEO, YMYL review, reputation/digital PR, AI-content governance, analytics, experiments, incident response, and Search Console/Bing/merchant/local monitoring.

Create portfolio governance for approximately 15 domains: distinct entity/audience/value requirements; allowed code/template reuse; content reuse/syndication rules; doorway and scaled-content prevention; cross-domain linking; brand/schema separation; ownership/approval; quality gates; and segmented reporting.

Include per-page, pre-launch, monthly, quarterly, migration, and portfolio checklists; a source registry with review date; P0-P3 rollout; and measurement from discovery through revenue. Validate Markdown, links, schema guidance, internal consistency, and git diff. Report changed files, tests, evidence limitations, and unresolved decisions. Never fabricate identities, dates, expertise, reviews, locations, metrics, citations, or outcomes.

Implementation Roadmap

Priority	Outcome	Representative work
P0	Remove material risk	Backdating; fake identities/evidence; misleading schema; quotas; doorway/scaled content; guarantees
P1	Build the foundation	Site charters; query maps; page contracts; indexation; governance; accessibility; evidence review
P2	Operationalize	Local/ecommerce/international/YMYL modules; automation; dashboards; incident and migration controls
P3	Learn and compound	Original research/tools; measured UX/content tests; optional AI publisher formats

Success is a system, not a promise

The playbook succeeds when teams consistently publish independently valuable, truthful, accessible, technically reliable experiences; detect risks early; learn from controlled evidence; and connect visibility to qualified business outcomes. Search and AI systems remain external and volatile, so the operating advantage is disciplined quality plus fast, measured adaptation.

Updated source playbook

playbooks/content-seo/SPEC.md in the webdev repository